

Devraj Bijpuria

Data Engineer

Patna, Bihar | dbijpuria@gmail.com | +91 7870668609

linkedin.com/in/devraj-bijpuria | github.com/DevrajBijpuria

Education

Vellore Institute of Technology
B.Tech, Computer Science Engineering

Sep 2023 – Aug 2027
CGPA: 8.58

Technical Skills

Programming	Python, SQL, C++
Data Engineering & ETL	ETL/ELT Pipeline Development, dbt, Data Warehousing, Database Design, Relational Databases, Snowflake SQL, Streamlined Data Ingestion with pandas
Cloud & AWS Data Analytics	AWS, Oracle Cloud Infrastructure (OCI), Cloud Computing Fundamentals
SQL & Databases	PostgreSQL, Joins, Window Functions, Relational DB Design, Data Manipulation in SQL
Data Analysis	Data Manipulation, Data Cleaning, Data Visualization, Exploratory Data Analysis, Importing Data in Python, APIs in Python, Statistics
Machine Learning & AI	Supervised & Unsupervised Learning, scikit-learn, XGBoost, Ensemble Methods, Hyperparameter Tuning, Tree-Based Models, Linear Classifiers, Isolation Forest

Projects

Spotify Data Pipeline — Serverless AWS ETL tinyurl.com/spotifypipeline Aug 2026

- Moved beyond one-off scripts to scheduled, unattended infrastructure that runs end-to-end without intervention.
- Fully serverless ETL pipeline that extracts Spotify listening data daily and exposes it for SQL-based analytics.
- 5 AWS services — CloudWatch-triggered Lambda extraction via the Spotify API; raw data in S3; S3-event Lambda for transformation; Glue Crawler inferring schema into the Data Catalog; Athena for SQL querying.
- Transform Lambda normalizes each raw pull into three tables — songs, albums and artists — deduping and cleaning dates with pandas via an OAuth refresh-token flow, then archives the processed file back to S3.

Real-Time Data Pipeline — Snowflake CDC tinyurl.com/RealTimeDataPipe Aug 2026

- Batch loads miss changes between runs and lose history — built for near-real-time ingestion with automatic change tracking and no manual steps.
- Near-real-time pipeline streaming synthetic customer data into Snowflake with Slowly Changing Dimensions (SCD1 current state, SCD2 full history).
- Python/Faker generates CSVs; Apache NiFi (Dockerized on EC2) delivers them to S3; Snowpipe auto-ingests into a staging table; Snowflake Streams capture changes and scheduled Tasks merge SCD1 and SCD2 every minute.
- Keeps two targets in sync automatically: a current-state table (SCD1, stored-procedure merge) and a full-history table with start/end timestamps and an is-current flag (SCD2, stream-plus-view merge).

Signal Desk — Zero-Cost News Intelligence Platform tinyurl.com/SignalDesk11 Aug 2026

- News apps surface headlines but not source credibility or emerging stories — set out to solve both at zero infrastructure cost.
- Serverless news platform across 4 verticals with credibility scoring, real-time sentiment tracking, and a newspaper-style reading experience.
- Netlify Scheduled Functions and Blobs for orchestration; source-legitimacy scoring; Tavily discovery search; Bluesky/YouTube sentiment aggregation; page-flip UI with StPageFlip.
- Rule-based legitimacy scoring (source-tier map plus a corroboration bump, with the reason shown on every story — no ML in the loop); a cron function sweeps, dedupes and scores four times a day into one Netlify Blob, so page loads read the cached sweep instead of hitting live feeds.

Achievements

- NextGen Hackathon (Finalist) — Built an ML-driven Intrusion Detection System prototype combining signature rules, an Isolation Forest anomaly detector and threat-intel lookups into a hybrid, real-time risk-scoring pipeline.
- NPTEL — Marketing Analytics.
- Data Engineering, Machine Learning Scientist and Data Analyst tracks — DataCamp.

Certifications

- **AWS Certified Cloud Practitioner — Amazon** Aug 2026
- **Generative AI Professional — Oracle** Jul 2025
- **Security, Compliance and Identity Fundamentals (SC-900) — Microsoft** Jun 2025